Machine Learning with TensorFLow Object Detection running on Docker

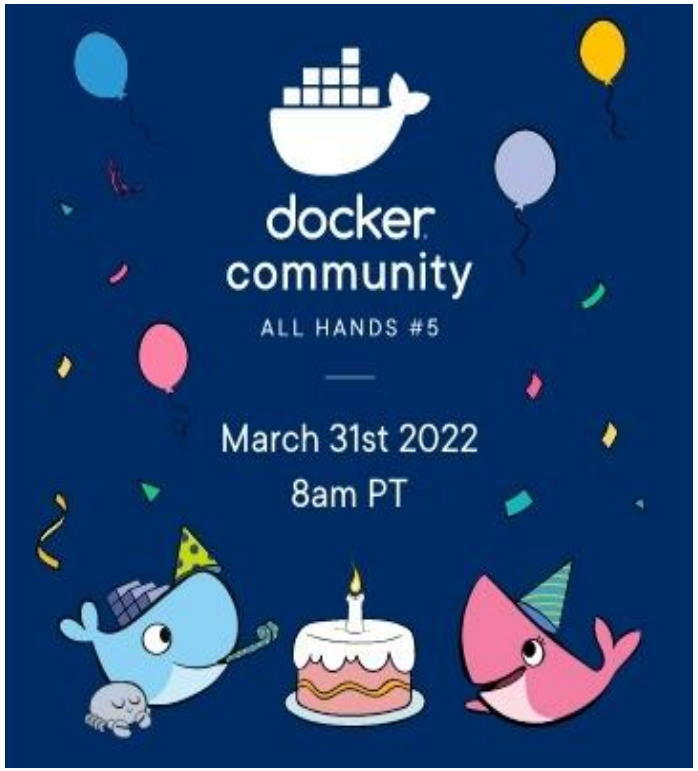

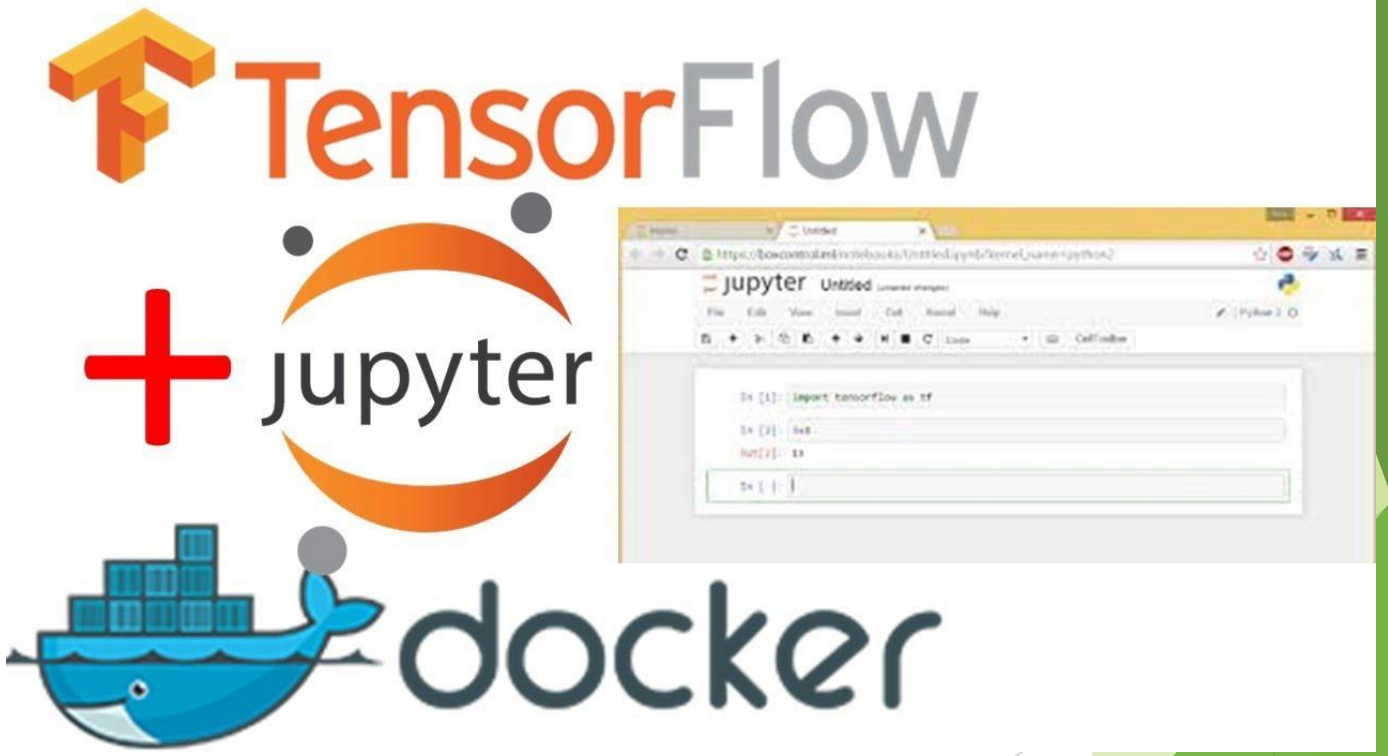

About the Speaker

Yeluri, Madhu Kumar
Principal Cloud Architect
@ T-Systems International
Budapest, Hungary

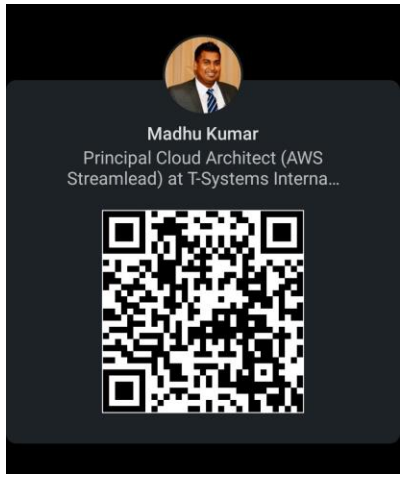

- Yeluri, Madhu Kumar

- Madhu is a qualified Principal Cloud Architect and Senior DevSecOps Consultant with over 21 years of IT experience working across multiple regions including Asia, Middle East, the US, Europe and the UK. He is helping many customers transform their business using the cloud. He is leading diverse teams to driving change and deliver business value at scale.

- A certified Amazon Web Services (AWS) Solution Architect and Security Specialist. Product Owner for Container services (Docker, K8s, AWS ECS and EKS). He has worked with many Cloud Partners/Providers (AWS, Rackspace, Wipro, Google, Oracle, Azure, IBM and Vodafone) and successfully managed and implemented multiple Cloud migration projects replacing business-critical core legacy systems across the Telecom, Financial, Banking, Insurance, Retail, and Government sectors.

Agenda

- Surprising Facts about Docker Adoption
- TensorFlow
- Create the Dockerfile
- Build the Dockerfile (Demo)
- Run the Docker Image (Demo)
- Object Detection API (Demo)
- Conclusion
- Questions

Surprising Facts about Docker Adoption

- **Nearly One Quarter of Companies Have Adopted Docker**
- **Docker Now Runs on More Than 20% of Hosts**
- **Docker Usage Rates Increase with Infrastructure Size**
- **Half of Docker Environments Are Orchestrated**
- **The Average Size of a Docker Deployment Has Grown 75% in One Year**
- **The Most Widely Used Images Are NGINX, Redis, and Postgres**
- **The Median Docker Organization Runs Eight Containers per Host**
- **Orchestrated Containers Churn 12x Faster**

Source: https://www.datadoghq.com/docker-adoption

TensorFlow

- **TensorFlow** is a free and open-source software library for machine learning and artificial intelligence. It can be used across a range of tasks but has a particular focus on training and inference of deep neural networks.

- TensorFlow was developed by the Google Brain team for internal Google use in research and production. The initial version was released under the Apache License 2.0 in 2015. Google released the updated version of TensorFlow, named TensorFlow 2.0, in September 2019.

Create the Dockerfile

- **FROM** "ubuntu:bionic"
- **RUN** apt-get update && yes | apt-get upgrade
- **RUN** mkdir -p /tensorflow/models
- **RUN** apt-get install -y git python-pip
- RUN pip install --upgrade pip
- **RUN** pip install tensorflow
- **RUN** apt-get install -y protobuf-compiler python-pil python-lxml
- **RUN** pip install jupyter
- RUN pip install matplotlib
- **RUN** git clone https://github.com/tensorflow/models.git /tensorflow/models
- **WORKDIR** /tensorflow/models/research
- **RUN** protoc object_detection/protos/*.proto --python_out=.
  RUN export PYTHONPATH=$PYTHONPATH:`pwd`:`pwd`/slim
- **RUN** jupyter notebook --generate-config --allow-root
- RUN echo "c.NotebookApp.password = u'sha1:6a3f528eec40:6e896b6e4828f525a6e20e5411cd1c8075d68619'" >> /root/.jupyter/jupyter_notebook_config.py
- **EXPOSE** 8888
- **CMD** ["jupyter", "notebook", "--allow-root", "--notebook-dir=/tensorflow/models/research/object_detection", "--ip=0.0.0.0", "--port=8888", "--no-browser"]

Demo

Build the Docker image

```
~/Downloads/docker-allhands5/tensorflow docker build -t tensorflow .
[+] Building 2.8s (19/19) FINISHED
 => [internal] load build definition from Dockerfile
 => => transferring dockerfile: 37B
 => [internal] load .dockerignore
 => => transferring context: 2B
 => [internal] load metadata for docker.io/library/ubuntu:bionic
 => [ 1/15] FROM docker.io/library/ubuntu:bionic@sha256:d8ac28b7bec51664c6b71a9dd1d8f788127ff310b
 => CACHED [ 2/15] RUN apt-get update && yes | apt-get upgrade
 => CACHED [ 3/15] RUN mkdir -p /tensorflow/models
 => CACHED [ 4/15] RUN apt-get install -y git python-pip
 => CACHED [ 5/15] RUN pip install --upgrade pip
 => CACHED [ 6/15] RUN pip install tensorflow
 => CACHED [ 7/15] RUN apt-get install -y protobuf-compiler python-pil python-lxml
 => CACHED [ 8/15] RUN pip install jupyter
 => CACHED [ 9/15] RUN pip install matplotlib
 => CACHED [10/15] RUN git clone https://github.com/tensorflow/models.git /tensorflow/models
 => CACHED [11/15] WORKDIR /tensorflow/models/research
 => CACHED [12/15] RUN protoc object_detection/protos/*.proto --python_out=.
 => CACHED [13/15] RUN export PYTHONPATH=$PYTHONPATH:`pwd`:`pwd`/slim
 => CACHED [14/15] RUN jupyter notebook --generate-config --allow-root
 => CACHED [15/15] RUN echo "c.NotebookApp.password = u'sha1:6a3f528eec40:6e896b6e4828f525a6e20e5
 => exporting to image
 => => exporting layers
 => => writing image sha256:e3850d4859c6ee479b610d13a32cbd882bade3d67beffa7f004fa8af5e78c5ea
 => => naming to docker.io/library/tensorflow

~/Downloads/docker-allhands5/tensorflow
```

- **Build the Docker image**
- docker build -t tensorflow .

Run the Docker Image

- **Run the Docker container**
  - docker run --rm --name tensorflow -p 8888:8888 -d tensorflow

- **Run the Docker container**
  - Open http://localhost:8888

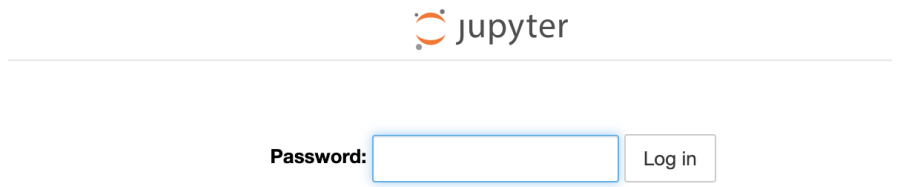

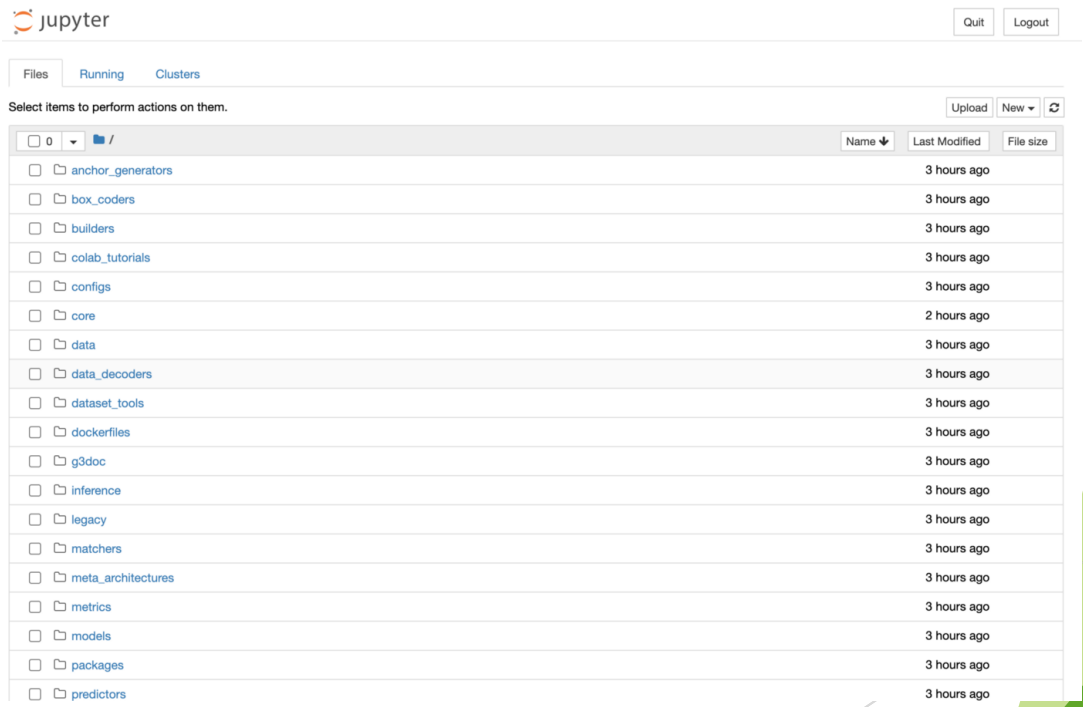

Object Detection API Demo

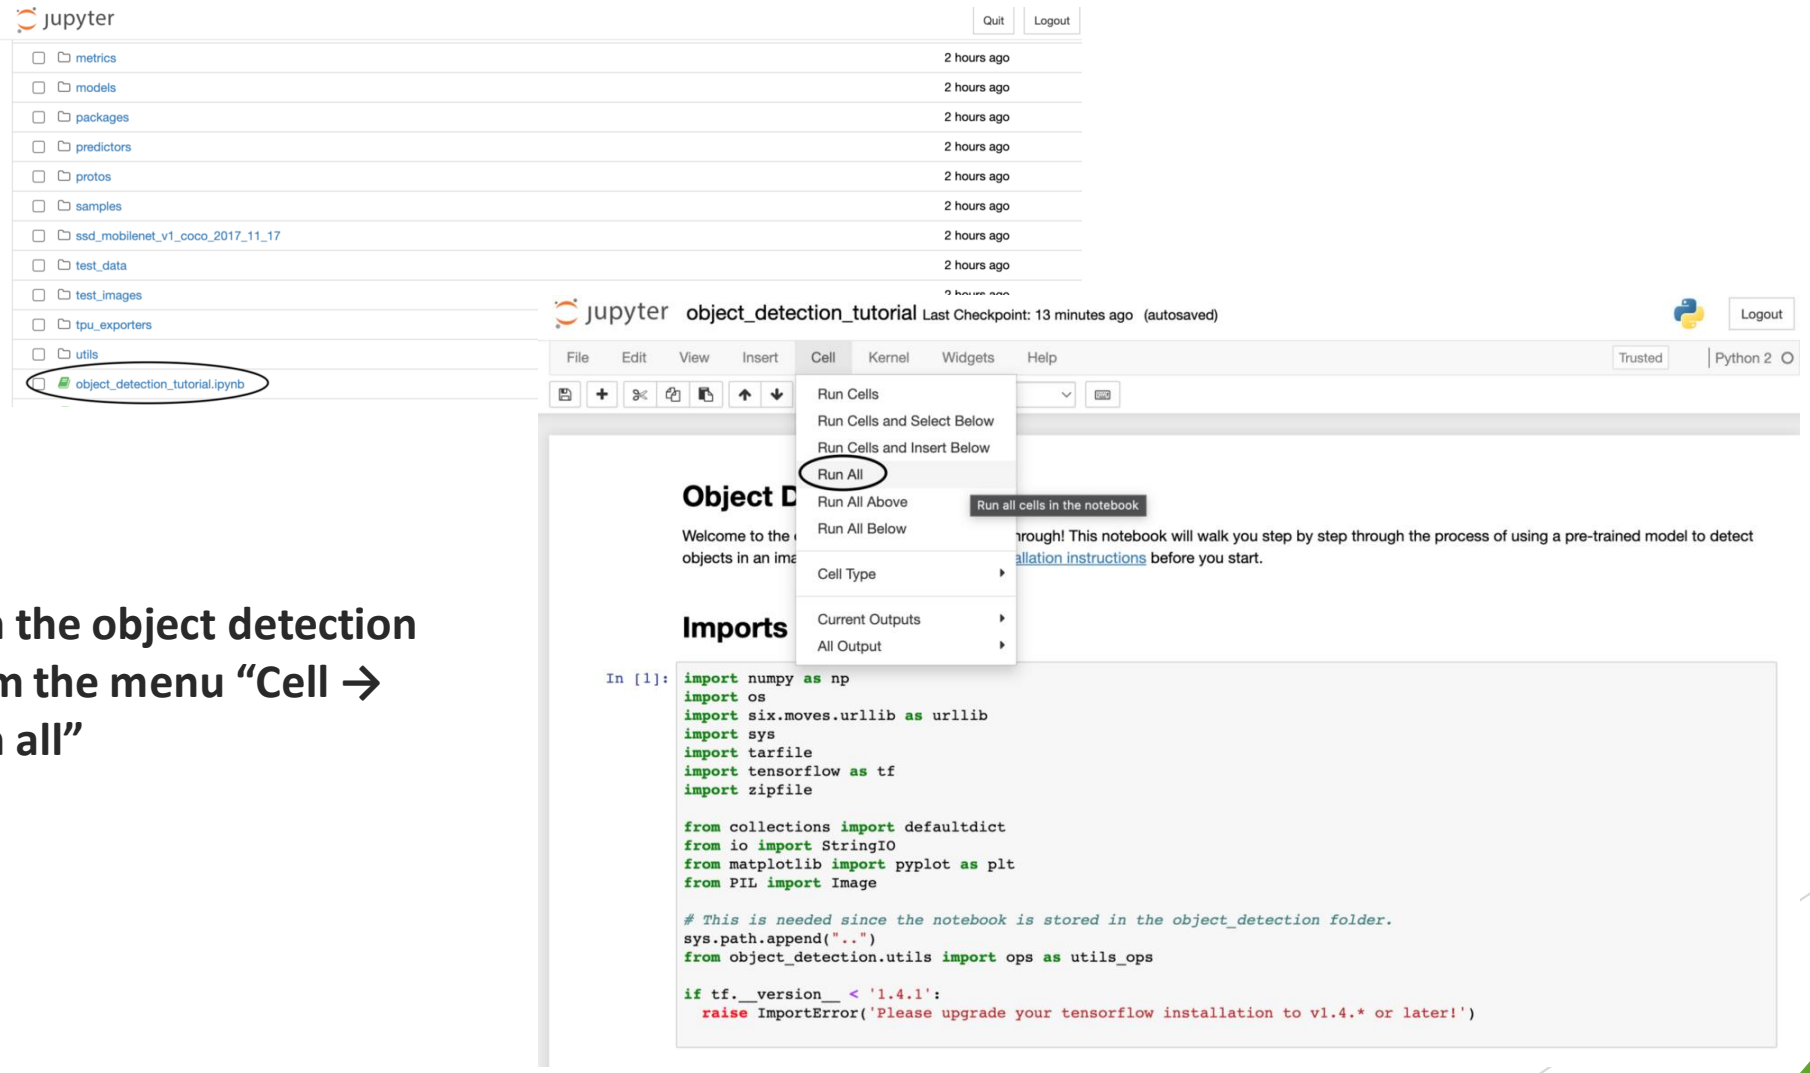

Run the object detection from the menu “Cell → Run all”

Results of the object detection

jupyter object_detection_tutorial Last Checkpoint: 6 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Run Code

```
    output_dict['detection_boxes'],
    output_dict['detection_classes'],
    output_dict['detection_scores'],
    category_index,
    instance_masks=output_dict.get('detection_masks'),
    use_normalized_coordinates=True,
    line_thickness=8)
plt.figure(figsize=IMAGE_SIZE)
plt.imshow(image_np)
```

jupyter object_detection_tutorial Last Checkpoint: 6 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

Run Code

```
    output_dict['detection_boxes'],
    output_dict['detection_classes'],
    output_dict['detection_scores'],
    category_index,
    instance_masks=output_dict.get('detection_masks'),
    use_normalized_coordinates=True,
    line_thickness=8)
plt.figure(figsize=IMAGE_SIZE)
plt.imshow(image_np)
```

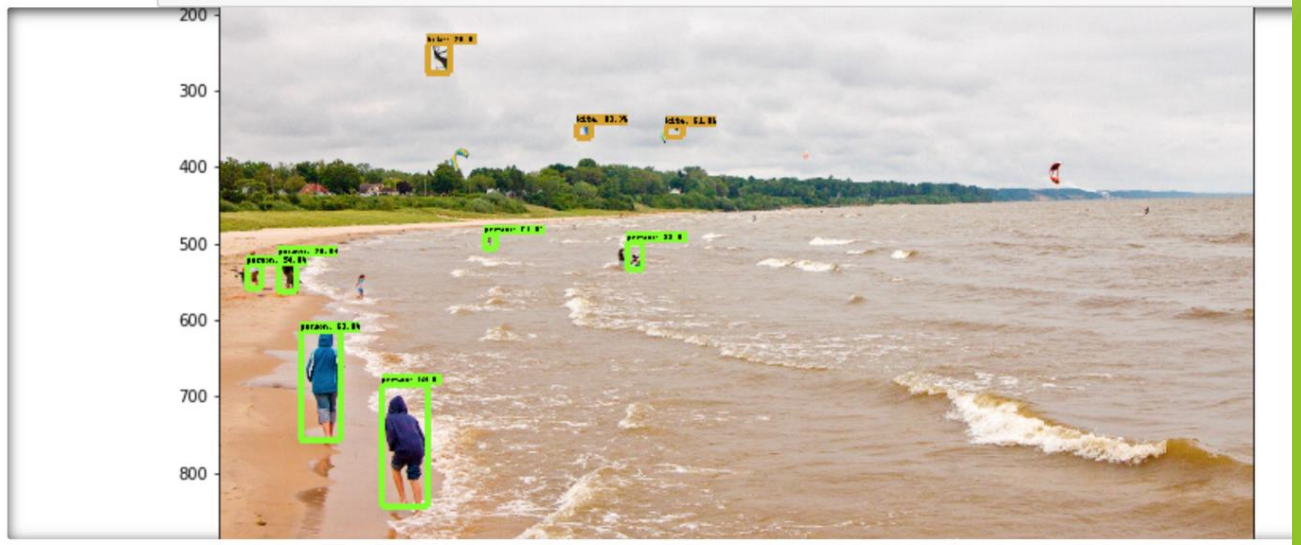

Stop the TensorFlow Docker container

```
docker rm -f tensorflow
```

Conclusion

Why Docker?

Docker is my favourite containerisation platform. Why? Docker provides a way to run applications securely isolated in a container, packaged with all its dependencies and libraries which are required for the application to run. Please refer to my other Docker blogs for more learning on Docker.

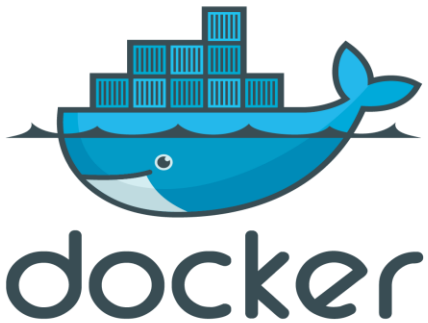

TensorFlow

Why TensorFlow?

TensorFlow supports various platforms for deploying ML models, be it desktop, mobile, web, or even cloud. Models can be trained using different programming languages like Python, JavaScript, or Swift.

As per the StackOverflow Developers Survey 2020, TensorFlow is one of the most popular frameworks among developers. Around 65% of the surveyed respondents have expressed their interest in continuing to develop models using TensorFlow. Also, with Google's support, the library will be enhanced regularly to fulfill the growing needs of developers.

Questions